

Exploring the unordered_map

In this exercise, we'll spend a little time understanding how the STL implements hash tables. We've already seen that the `unordered_map` uses something like chaining (though whether the "buckets" are "chains" isn't totally clear from the documentation). But how many buckets are there? How big are they? And so on.

Understanding the Basics

Write a small program as follows:

Create a "basic" `unordered_map` where `key_type` and `mapped_type` are both `int`. (That is, we'll just store ints in the hash table for now.)

Use the `bucket_count` and `load_factor` methods to learn about the default initial state of an `unordered_map`.

Add 10 unique values to the map. After each addition, check the `bucket_count` and `load_factor`. Run the program, and based on these results, make a hypothesis about how the map is resized. According to your hypothesis, what should the `bucket_count` and `load_factor` be after (say) 100 values have been added to the map?

Now add to your program—add 90 more values, so there are a total of 100 values in the map. Output the `bucket_count` and `load_factor`; also use `bucket_size` to find out the max number of elements stored in any one bucket. How does that number relate to the `load_factor`?

The Danger of Bad Hash

OK, let's implement the `Point` class, along with a bad hash function; let's see what that does to our maps. The `Point` class can be super-simple:

```
class Point {
private:
    int x, y;
public:
    Point(int x, int y) : x(x), y(y) {};
    bool operator==(const Point& p) const {return x == p.x && y == p.y;}
    int getX() const {return x;}
    int get&() const {return y;}
}

std::ostream& operator<< (std::ostream& out, const Point& p) {
    out << "(" << p.getX() << ", " << p.getY() << ")";
}
```


But we also need to implement a hash function that will work with `unordered_map`. To do that, we have to specialize the `std::hash` template to work for `Point`. We'll use the bad hash function `x+y`; to see how to incorporate that, look at the example code at [http://en.cppreference.com/w/cpp/utility/hash/operator\(\)](http://en.cppreference.com/w/cpp/utility/hash/operator()).

Now write a main program that creates a hash table for `Points`. Add 100 points to the table, from (0,0) to (9,9). What are the `bucket_count` and `load_factor` for this table? Are they the same as for the previous table of 100 ints? Does the “bad hash function” seem to be affecting these values?

Follow the example at http://www.cplusplus.com/reference/unordered_map/unordered_map/bucket_count/ to print the contents of each bucket. Does the “bad hash function” seem to be affecting these results?